

MapReduce

Databases Three

Otago Polytechnic
Dunedin, New Zealand

Problem: counting splatts

Suppose we want to count the number of splatts in our collection.

- Relational : `select count(*) from splatts;`
- Document : If we had a splatts collection, we could do this - `db.splatts.count()` - but we don't have one.
- We need to tell MongoDB how to find and count our splatts.

The data

```
{
  "_id" : ObjectId("5416717562696259b8010000"),
  "email" : "imp@casterlyrock.com",
  "follower_ids" : [ ObjectId("5416717562696259b8000000") ],
  "name" : "Tyrion Lannister",
  "splatts" : [
    {
      "_id" : ObjectId("5416727962696259b8030000"),
      "body" : "user 2 splatt 2",
      "created_at" : ISODate("2014-09-15T05:00:41.376Z")
    }
  ]
}
```

The plan

- 1 Iterate over the users¹;
- 2 Count the splatts belonging to each one;
- 3 Add up the counts.

This is a common pattern, and it has a name: MapReduce.

¹Our data set is small, but this could be parallelised over a large data set.

MapReduce

- 1 Map: For each user, count the splatts.
- 2 Reduce: Sum up the individual counts.

MapReduce implementation

We need to supply JavaScript code to MongoDB telling it how to perform the Map and Reduce steps

- We write a map function that is called one time on each user in the collection. It will emit the count of each user's splatts.
- We write a reduce function that sums up the output from the map function.
- We supply both of these functions to MongoDB's `mapReduce` command.

Map

```
var map = function() {  
  var length = 0;  
  if(this.splatts) {  
    length = this.splatts.length  
  }  
  emit ("count", length);  
};
```

Map results

After applying the map to a collection with 3 users, the result set may look like this:

```
results = [  
  {key: "count", value: 4 },  
  {key: "count" value: 3 },  
  {key: "count", value: 3 }  
]
```


Reduce

```
var reduce = function(key, val) {  
  var data = 0;  
  val.forEach(function(v) {  
    data += v;  
  });  
  return data;  
}
```

This reduce function will be called one time for each key value in the map results.

MapReduce

We use our map and reduce functions in a call to MongoDB's mapReduce command.

```
db.users.mapReduce(  
  map,  
  reduce,  
  {  
    out: {inline: 1}  
  }  
)
```

The final result

```
{
  "results" : [
    {
      "_id" : "count",
      "value" : 10
    }
  ],
  "timeMillis" : 0,
  "counts" : {
    "input" : 3,
    "emit" : 3,
    "reduce" : 1,
    "output" : 1
  },
  "ok" : 1,
}
```

In Ruby

The Mongoid library provides access to MongoDB's MapReduce functionality:

```
map = %Q{ function() {  
    var length = 0;  
    if(this.splatts) {  
        length = this.splatts.length  
    }  
    emit ("count", length);  
}  
}  
  
reduce = %Q{ function(key, val) {  
    var data = 0;  
    val.forEach(function(v) {  
        data += v;  
    })  
    return data;  
}  
}
```

In Ruby

```
User.map_reduce(map,reduce).out(inline: true)
```

Today's lab

- 1 Adapt the test scripts you wrote earlier this semester to enter some sample data into your MongoDB database.
- 2 Use MapReduce from within the MongoDB shell to get a count of the number of splatts.
- 3 Add a function to your Rails UsersController to return the total number of splatts.